

Erlang Bootcamp 2

Erlang Bootcamp 2:

Agenda

- (Re)Introducing gen_server
- Supervisors
- OTP Applications
- rebar

Agenda

- TCP Sockets
- Distributed Erlang
- Webmachine

Introducing `gen_server`

Introduction

- Abstracts away all low-level server details
- Integrates with the rest of OTP
- Makes writing servers a snap!

What is a behavior?

- Defines an API
- Contract between modules
- Provide support functions (optional)

gen_server

- Implemented as a behavior
- Defines an API for server logic module
- Provides server logic support functions

Startup & Shutdown

init/ I

- Called during startup
- Responsible for generating initial server state
- Returns {ok, State} or {stop, Reason}

`init([]) ->`

`{ok, #state{}}.`

terminate/2

- Called during teardown
- Responsible for cleaning up resources
- Return value is ignored

terminate(_Reason, _State) ->


ok.

Message Types

‘Call’ Messages

- Client is blocked until server replies
- Default timeout interval is 5 seconds
- Clients use `gen_server:call/3,4`


'Call' Messages



‘Cast’ Messages

- Asynchronous from the client’s POV
- Server isn’t required to reply
- Clients use `gen_server:cast/2`


'Cast' Messages



‘Info’ Messages

- Asynchronous
- Server isn’t required to reply
- Delivered via “raw” message send

'Info' Messages



Handling Messages

‘Call’ Messages

```
handle_call(current_time, _From, State) ->  
    {reply,  
        calendar:now_to_local_time(erlang:now()),  
        State};
```

‘Cast’ Messages

```
handle_cast(foo, State) ->  
  do_something(foo),  
  {noreply, State};
```

‘Info’ Messages

```
handle_info(reset, State) ->  
    {noreply, reset(State)};
```


Lab Time!

- Implement an in-memory cache using `gen_server`. Cached values should be stored in an Erlang dictionary.
- Cached values should be stored in a tuple containing the value and a timestamp. Values older than X seconds should be deleted.
- The `gen_server` should expose 4 client functions: `get/1`, `put/2`, `del/1`, `count/0`.

Supervisors

What are supervisors?

- Processes which monitor other processes
- Supervisors can supervise other supervisors
- Very configurable (good and bad)

How Do They Work?

- Each child process is linked to its supervisor
- Supervisors restart crashed children

Configuring Children

{RestartPolicy, Children}

Restart Policy

- Controls how a crashed child is handled
- `one_for_one`, `one_for_all`, `rest_for_one`
- Applied to all children

Restart Interval

- Describes the “fatal crash scenario”
- Prevents supervisor from restarting children forever

{one_for_one, 10, 60}

Child Specification

- How to start the child
- How to manage the child
- Which code belongs to the child

{Id, StartFun, Restart, Shutdown, Type, Modules}

Start Function

- Child entry point
- MFA-formatted tuple: {Module, Fun, Args}
- Starts and links the child to the supervisor

{time_server,

{time_server, start_link, []}

Restart

- Determines when the restart policy is used
 - Permanent: always restarted
 - Temporary: never restarted
 - Transient: restarted when it crashes

```
{time_server,  
  {time_server, start_link, []},  
  permanent
```

Shutdown

- How to stop the child
 - `brutal_kill`: shoot it in the head (`kill -9`)
 - `<integer>`: wait for *X* milliseconds before `brutal_kill`-ing
 - `infinity`: wait forever

```
{time_server,  
  {time_server, start_link, []},  
  permanent, 2000
```


Type

- worker: Normal child process
- supervisor: Nested supervisor. Should use 'infinity' shutdown

```
{time_server,  
  {time_server, start_link, []},  
  permanent, 2000, worker
```

Modules

- Modules which make up the child process
- Normally a single element list

```
{time_server,  
  {time_server, start_link, []},  
  permanent, 2000, worker,  
  [time_server]}
```

Lab Time!

- Write a supervisor for the cache server. Make sure the cache server starts successfully after you call the supervisor's `start_link()`.

Working with OTP Applications


What is an OTP Application?

- Recommended format for organizing projects
- Compatible with the OTP libraries and the Erlang runtime
- Compatible with tools like rebar, proper, and dialyzer

OTP App Structure

- Top-level directory named after the app
- Erlang source kept in 'src'
- Compiled `.beam` files kept in 'ebin'
- Header files kept in 'include'
- App descriptor (`.app`) stored in 'ebin'

OTP App Structure



Application Descriptor

- Plain text file named after the application
- Describes app metadata:
 - Entry point
 - Version
 - Included modules
 - Registered names
- Ends in `' .app'`

.app Example

```
{application, erlz,  
  [  
    {description, "Wrapper for the fastlz compression library"},  
    {vsn, "0.1"},  
    {registered, []},  
    {applications, [kernel,  
                    stdlib]},  
    {env, []}  
  ]}.  
}
```

Minimum .app Requirements

- Application
- Version
- Applications list
- Modules list

.app Example

```
{application, webmachine,  
  [{description, "webmachine"},  
   {vsn, "1.8.0"},  
   {modules, [webmachine,  
               webmachine_app,  
               webmachine_decision_core,  
               [...] ]},  
  ]},  
 {registered, []},  
 {mod, {webmachine_app, []}},  
 {env, [  
   {dispatch_list, []}  
 ]},  
 {applications, [kernel, stdlib, crypto, ssl]}}.
```

.app Example

```
{application, webmachine,  
  [{description, "webmachine"},  
   {vsn, "1.8.0"},  
   {modules, [webmachine,  
               webmachine_app,  
               webmachine_decision_core,  
               [...] ]},  
  ]},  
{registered, []},  
{mod, {webmachine_app, []}},  
{env, [  
    {dispatch_list, []}  
  ]},  
{applications, [kernel, stdlib, crypto, ssl]}]}
```

.app Example

```
{application, webmachine,  
  [{description, "webmachine"},  
   {vsn, "1.8.0"},  
   {modules, [webmachine,  
               webmachine_app,  
               webmachine_decision_core,  
               [...] ]},  
  ]},  
{registered, []},  
{mod, {webmachine_app, []}},  
{env, [  
    {dispatch_list, []}  
  ]},  
{applications, [kernel, stdlib, crypto, ssl]}}.
```

Application Behavior

- Abstracts the ‘entry point’ concept
- Responsible for starting the supervision tree
- Two functions: `start/2` and `stop/1`


```
-module (foo_app) .  
  
-behaviour (application) .  
  
%% Application callbacks  
-export ([start/2, stop/1]) .  
  
start (_StartType, _StartArgs) ->  
    foo_sup:start_link() .  
  
stop (_State) ->  
    ok.
```

Pains of Building an OTP App

- Maintaining the modules list
- Building source
 - Include paths
 - Dependencies
- Automating builds



rebar

What is rebar?

- Project generator
- Build tool
- Dependency manager

rebar as Project Generator

Project Generator

```
rebar create-app appid=app_name
```

- Generates OTP dirs (ebin, src)
- Creates top-level supervisor and app modules
- Creates `app_name.app`

rebar as Build Tool

Supported File Types

- Erlang source (`.erl`)
- Protocol Buffers (`.proto`)
- Linked-in drivers (`.c`, `.cpp`, `Makefiles`)
- NIFs

Supported File Types

- Leex/Yecc grammars (`.xrl`, `.yrl`)
- Eunit (tests & fixtures)
- Common Test
- QuickCheck

Controlling The Build Process

- `rebar.config` controls **all**
- Kept in top of application directory

Example

rebar.config

```
{erl_opts, [debug_info, warnings_as_failures]}.
```

```
{deps, [{erlzmq2, ".*",  
           {git, "git://github.com/zeromq/erlzmq2.git",  
              "HEAD"}}]}.
```

Example

rebar.config

```
{erl_opts, [debug_info,  
             warnings_as_failures]}.  
  
{deps, [{erlzmq2, ".*",  
            {git, "git://github.com/zeromq/erlzmq2.git",  
              "HEAD"}}]}.
```

Compiling Erlang

- Sanity checks `.app` file on each compile
- Use `erl_opts` to configure compiler
- Automagically populate modules list (!!!)

rebar: `rebar compile`

rebar as Dependency Manager

rebar

Dependencies

- **Not** a package manager
- Dependencies on 3rd party *Erlang* source projects
- Supports `svn`, `git`, and `hg`

Declaring Dependencies

App Name App Version

↓ ↓

```
{deps, [{erlzmq2, ".*",  
               {git, "git://github.com/zeromq/erlzmq2.git",  
                  "HEAD"}]}].
```

↑ ↑

VCS Type VCS URL

↑

VCS Version/Branch/Tag

```
graph TD
    AN[App Name] --> erlzmq2
    AV[App Version] --> stars[.*]
    VT[VCS Type] --> git
    VCU[VCS URL] --> url["git://github.com/zeromq/erlzmq2.git"]
    VCBT[VCS Version/Branch/Tag] --> HEAD
```


Using Dependencies

- Stored in deps/
- Fetched when needed
- Automatically added to the build-time codepath

rebar: `rebar [check|get|
delete]-deps`

Lab Time!

- Package the cache server as an OTP application:
 - .app file
 - application behavior
- Build the project using rebar
- Verify the cache server starts after calling `'application:start(cache_server).'`

Network Programming with TCP Sockets

TCP Sockets

- Two categories: clients & servers
- Clients connect to servers
- Servers accept incoming connections
- Both use `gen_tcp` and `inet` modules to interact with sockets

gen_tcp

- Interface for working with TCP sockets
- Used by clients and servers
- Works with IPv4 and IPv6

gen_tcp Clients

- Use `connect/3,4` to open connection
- Use `send/2` and `recv/2,3` to communicate

```
{ok, Sock} = gen_tcp:connect("localhost",  
                             5432,  
                             [binary,  
                             {active, false}]).
```

```
gen_tcp:send(Sock, <<"Hello">>).
```

```
{ok, Sock} = gen_tcp:connect("localhost",  
                             5432,  
                             [binary,  
                             {active, false}]).
```

```
gen_tcp:send(Sock, <<"Hello">>).
```


Socket Options

- Initial values set during socket creation
- Additional values set via `inet:setopts/2`

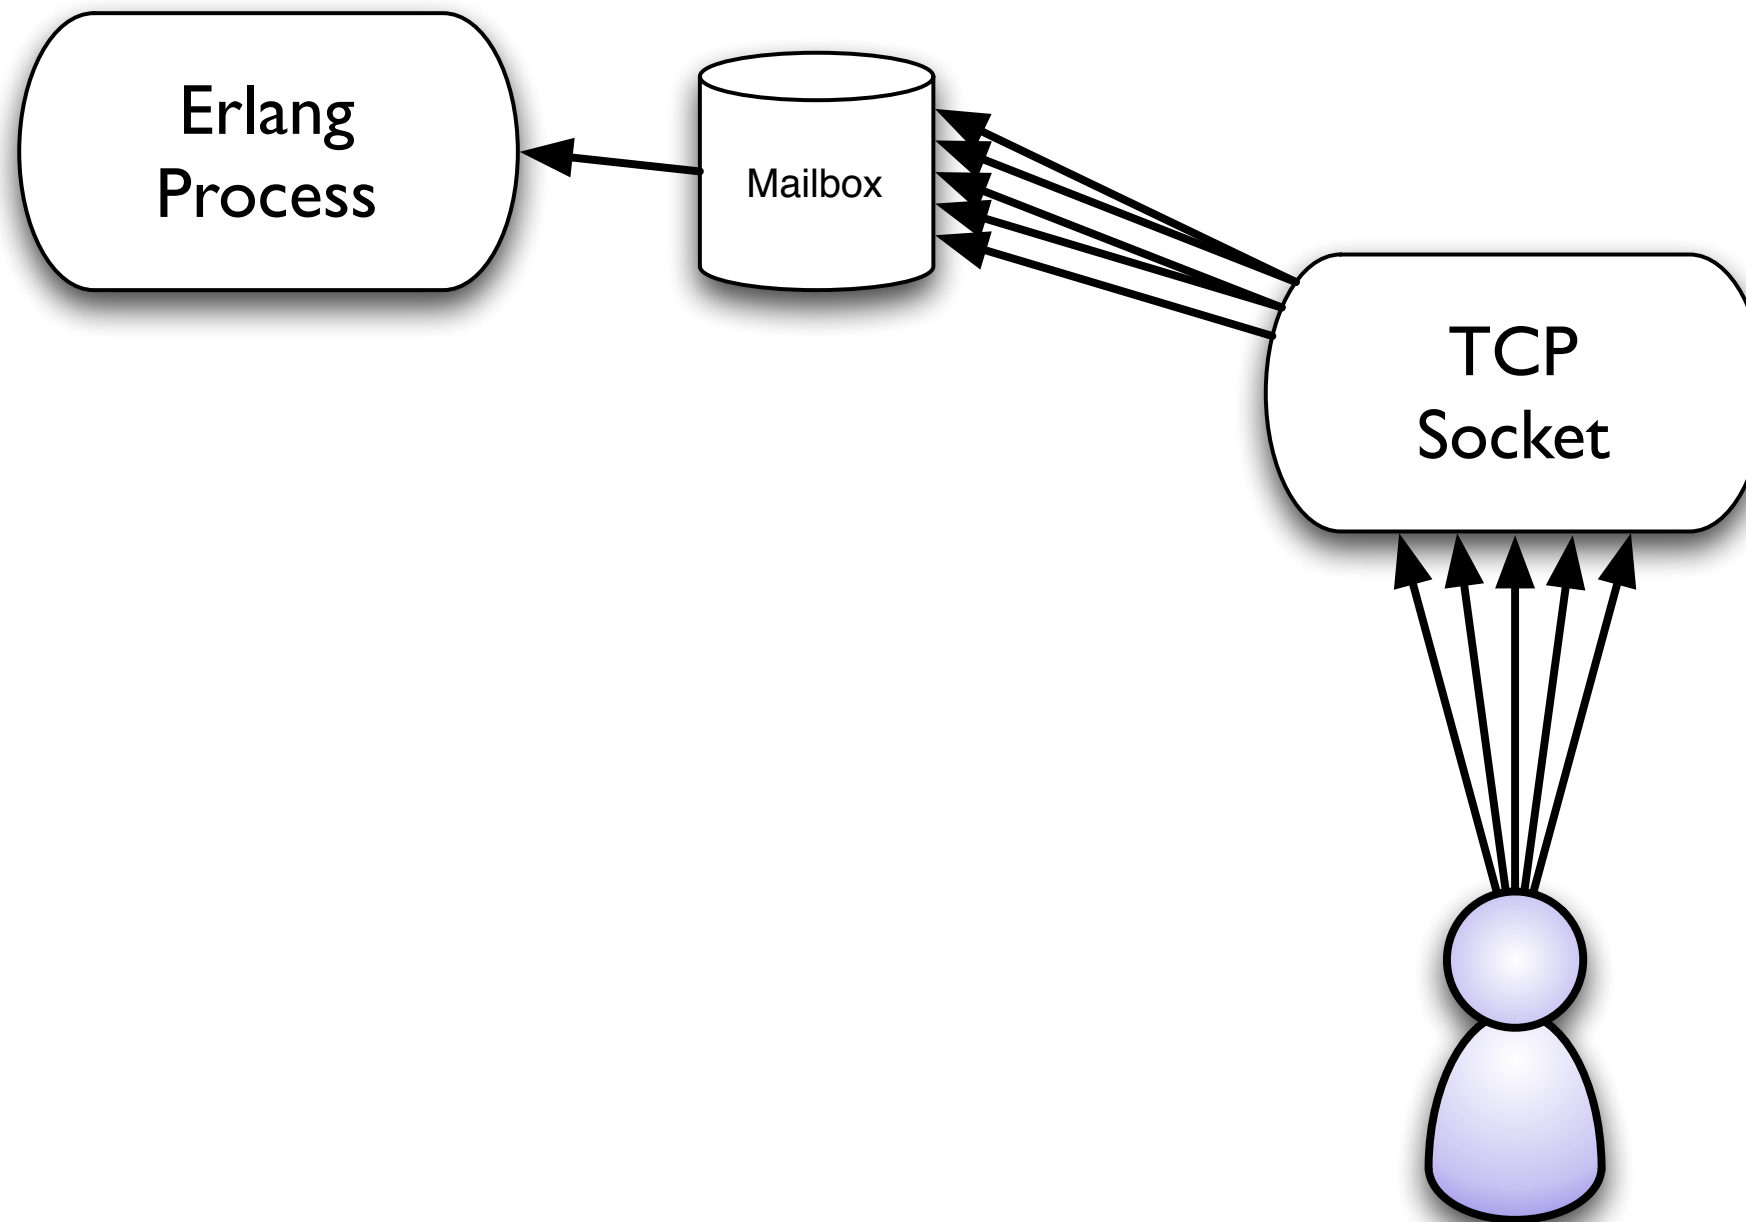
Socket Options

Option	Description
<code>{active, [true false once]}</code>	Flow control
<code>{packet, [raw N line binary http ...]}</code>	Enables data parsing
<code>{recvbuf, N}</code> <code>{sndbuf, N}</code>	Sets the size of the receive and send buffers

Flow Control

- Active - Erlang sends data as it's **received**
(Fast but prone to crashes)


{active, true}



Flow Control

- Active - Erlang sends data as it's **received**
(Fast but prone to crashes)
- Active Once - Erlang sends data as it's **acked**
(Slower but more stable)

{active, once}



Server Sockets

- Use `gen_tcp:listen/2` to open server socket
- Use `gen_tcp:accept/1, 2` in a tail recursive process
- Spawn a new process for each connection

Why gen_nb_server?

- Tired of writing accept loops
- Other libs were too heavy
- Push accept down into `prim_inet`

gen_nb_server

- `gen_server + sockets = fun`
- Listen on multiple IPs and ports
- Add/remove IPs and ports at runtime

Using gen_nb_server

- Implement all the gen_server callbacks
- `init/2` instead of `init/1`
- `sock_opts/0`
- `new_connection/4`

Socket Handling Tips

- Use active mode sockets whenever possible
- Use {active, once} to protect against flooding
- Use appropriate packet type to pre-parse incoming data

Supervisors As 'Process Factories'

Aren't Factories For Java?

NO.

Supervision Is A Good Idea

- Single point of creation
- Easier to trace execution path
- Supervision tree is easier to understand
- Makes debugging easier

simple_one_for_one

- Supervisor for a single process type (module)
- All children are started dynamically
- Crashed children are **not** restarted

```
-behaviour(supervisor).
```

```
%% API
```

```
-export([start_link/0,  
        new_handler/1]).
```

```
%% Supervisor callbacks
```

```
-export([init/1]).
```

```
-define(SERVER, ?MODULE).
```

```
start_link() ->
```

```
    supervisor:start_link({local, ?SERVER}, ?MODULE, []).
```

```
new_handler(Conn) ->
```

```
    supervisor:start_child(?SERVER, [Conn]).
```

```
init([]) ->
```

```
    {ok, {{simple_one_for_one, 0, 1},  
         [{conn_handler, {conn_handler, start_link, []}},  
          temporary, brutal_kill, worker, [conn_handler]}}}.
```



```
-behaviour(supervisor).
```

```
%% API
```

```
-export([start_link/0,  
        new_handler/1]).
```

```
%% Supervisor callbacks
```

```
-export([init/1]).
```

```
-define(SERVER, ?MODULE).
```

```
start_link() ->
```

```
    supervisor:start_link({local, ?SERVER}, ?MODULE, []).
```

```
new_handler(Conn) ->
```

```
    supervisor:start_child(?SERVER, [Conn]).
```

```
init([]) ->
```

```
    {ok, {{simple_one_for_one, 0, 1},  
         [{conn_handler, {conn_handler, start_link, []}},  
          temporary, brutal_kill, worker,  
          [conn_handler]]}}.
```

```
-behaviour(supervisor).
```

```
%% API
```

```
-export([start_link/0,  
        new_handler/1]).
```

```
%% Supervisor callbacks
```

```
-export([init/1]).
```

```
-define(SERVER, ?MODULE).
```

```
start_link() ->
```

```
    supervisor:start_link({local, ?SERVER}, ?MODULE, []).
```

```
new_handler(Conn) ->
```

```
    supervisor:start_child(?SERVER, [Conn]).
```

```
init([]) ->
```

```
    {ok, {{simple_one_for_one, 0, 1},
```

```
        [{conn_handler, {conn_handler, start_link, []}},  
         temporary, brutal_kill, worker, [conn_handler]}}}
```

```
-behaviour(supervisor).
```

```
%% API
```

```
-export([start_link/0,  
        new_handler/1]).
```

```
%% Supervisor callbacks
```

```
-export([init/1]).
```

```
-define(SERVER, ?MODULE).
```

```
start_link() ->
```

```
    supervisor:start_link({local, ?SERVER}, ?MODULE, []).
```

```
new_handler(Conn) ->
```

```
    supervisor:start_child(?SERVER, [Conn]).
```

```
init([]) ->
```

```
    {ok, {{simple_one_for_one, 0, 1},  
         [{conn_handler, {conn_handler, start_link, []}},  
          temporary, brutal_kill, worker, [conn_handler]}}}.
```

Lab Time!

- Write a socket front-end to the cache server using `gen_nb_server`.
- Each connection will be handled by a separate Erlang process. This means you will need to create a new process when `new_connection/4` is called.
- Clients will use a simple line-based protocol to send commands to the server:

```
get <key> \r\n
```

```
del <key> \r\n
```

```
put <key> <value> \r\n
```

- Servers responses can be one of three types:

```
+ok\r\n (no payload)
```


```
+ok <value>\r\n (with payload)
```

```
-err\r\n (error or not found)
```

Hint: *Socket options* `[list, {packet, line}, {active, once}]`
and the `string` module will be very helpful.

Distributed Erlang


Distributed Erlang




Distributed Erlang

- Each node maintains its own node list
- Nodes discover each other via
`net_adm:ping/1` or
`net_kernel:connect/1`
- New nodes “inherit” the node list of the node they ping


Node Discovery




Node Discovery




Node Discovery



Node Discovery



Dropping Nodes



Fully Meshed Networks Suck

Nodes	Connections
4	6
16	120
32	496 (!!!)

Registered Processes

- Associates a human-friendly name to a pid
- Default scope is local to the hosting node
- Use global name registry to register names across nodes

Note: Global names can be a bottleneck

Process Groups

- Named groups of processes across nodes
- Implement distributed process pools
- Use pg2 instead of pg

Note: process groups can cause problems in clusters experiencing membership churn.

Useful pg2 Functions

- `pg2:create/1`
- `pg2:join/2`
- `pg2:leave/1`
- `pg2:get_members/1`,
`pg2:get_closest_member/1`

Lab Time!

Goal: Create a new server which will round-robin between cache servers running on different nodes. All adds and deletes must be replicated to all cache servers.

The load balancing server will:

- Use `gen_server`
- Provide the same functional interface but delegate all work to the backend cache servers
- Register itself globally
- Provide a function backend servers can use to register themselves
- Use process monitors to remove dead backend servers

The backend cache servers will:

- Register themselves with the load balancer when they start. It is OK to hardcode the name.

webmachine

What is webmachine?

~~Framework~~

Toolkit

A toolkit for building RESTful HTTP resources

```
-module(hello_world_resource).  
-export([init/1, to_html/2]).  
  
-include_lib("webmachine/include/webmachine.hrl").  
  
init([]) ->  
    {ok, "Hello, world"}.  
  
to_html(Req, State) ->  
    Output = io_lib:format("<html><body>~s</body></html>",  
                           [State]),  
    {Output, Req, State}.
```

```
-module(hello_world_resource).  
-export([init/1, to_html/2]).  
  
-include_lib("webmachine/include/webmachine.hrl").  
  
init([]) ->  
    {ok, "Hello, world"}.  
  
to_html(Req, State) ->  
    Output = io_lib:format("<html><body>~s</body></html>",  
                            [State]),  
    {Output, Req, State}.
```

```
-module(hello_world_resource).  
-export([init/1, to_html/2]).  
  
-include_lib("webmachine/include/webmachine.hrl").  
  
init([]) ->  
    {ok, "Hello, world"}.  
  
to_html(Req, State) ->  
    Output = io_lib:format("<html><body>~s</body></html>",  
                           [State]),  
    {Output, Req, State}.
```



```
-module(hello_world_resource).  
-export([init/1, to_html/2]).  
  
-include_lib("webmachine/include/webmachine.hrl").  
  
init([]) ->  
    {ok, "Hello, world"}.  
  
to_html(Req, State) ->  
    Output = io_lib:format("<html><body>~s</body></html>",  
                           [State]),  
    {Output, Req, State}.
```

```
-module(hello_world_resource).  
-export([init/1, to_html/2]).  
  
-include_lib("webmachine/include/webmachine.hrl").  
  
init([]) ->  
    {ok, "Hello, world"}.  
  
to_html(Req, State) ->  
    Output = io_lib:format("<html><body>~s</body></html>",  
                                [State]),  
    {Output, Req, State}.
```

```
GET /test HTTP/1.1
Host: localhost:8000
User-Agent: Mozilla/5.0
Accept: text/html;q=0.9,*/*;q=0.8
Accept-Language: en-us,en;q=0.5
Accept-Encoding: gzip,deflate
Accept-Charset: UTF-8,*
Keep-Alive: 300
Connection: keep-alive

HTTP/1.x 200 OK
Server: MochiWeb/1.1 WebMachine/1.3
Date: Mon, 22 Jun 2009 02:27:46 GMT
Content-Type: text/html
Content-Length: 78
```

```
-module(hello_world_resource).  
-export([init/1, to_html/2]).  
-export([generate_etag/2]).  
  
-include_lib("webmachine/include/webmachine.hrl").  
  
init([]) ->  
    {ok, "Hello, world"}.  
  
to_html(Req, State) ->  
    Output = io_lib:format("<html><body>~s</body></html>",  
                           [State]),  
    {Output, Req, State}.  
  
generate_etag(Req, State) ->  
    {mochihex:to_hex(crypto:md5(State)), Req, State}.
```

GET /test HTTP/1.1
Host: localhost:8000
User-Agent: Mozilla/5.0
Accept: text/html;q=0.9,*/*;q=0.8
Accept-Language: en-us,en;q=0.5
Accept-Encoding: gzip,deflate
Accept-Charset: UTF-8,*
Keep-Alive: 300
Connection: keep-alive

HTTP/1.x 200 OK
Server: MochiWeb/1.1 WebMachine/1.3
Etag: bc6e6f16b8a077ef5fbc8d59d0b931b9
Date: Mon, 22 Jun 2009 02:29:46 GMT
Content-Type: text/html
Content-Length: 78

GET /test HTTP/1.1

Host: localhost:8000

.

.

.

If-None-Match: bc6e6f16b8a077ef5fbc8d59d0b931b9

HTTP/1.x 304 Not Modified

Server: MochiWeb/1.1 WebMachine/1.3

Etag: bc6e6f16b8a077ef5fbc8d59d0b931b9

Date: Mon, 22 Jun 2009 02:30:00 GMT

```
-module(hello_world_resource).  
-export([init/1, to_html/2]).  
-export([generate_etag/2, encodings_provided/2]).  
  
-include_lib("webmachine/include/webmachine.hrl").  
  
init([]) ->  
    {ok, "Hello, world"}.  
  
to_html(Req, State) ->  
    Output = io_lib:format("<html><body>~s</body></html>",  
                           [State]),  
    {Out, Req, State}.  
  
generate_etag(Req, State) ->  
    {mochihex:to_hex(crypto:md5(State)), Req, State}.  
  
encodings_provided(Req, State) ->  
    {[{ "gzip", fun(X) -> zlib:gzip(X) end}], Req, State}.
```

GET /test HTTP/1.1
Host: localhost:8000

.
.
.

Accept-Encoding: gzip, deflate

HTTP/1.x 200 OK
Server: MochiWeb/1.1 WebMachine/1.3
Etag: bc6e6f16b8a077ef5fbc8d59d0b931b9
Date: Mon, 22 Jun 2009 02:46:57 GMT
Content-Type: text/html
Content-Length: 71
Content-Encoding: gzip

Request Routing

http://foo.com/items

`{"items"}, grocery_item_resource, []}.`



URL path



Resource module



Init params

http://foo.com/items/chocolate

`{"items", item}, grocery_item_resource, []}.`



URL path



URL variable



Resource module



Init params

GET

```
-module(hello_world_resource).  
-export([init/1,to_html/2]).  
  
-include_lib("webmachine/include/webmachine.hrl").  
  
init([]) ->  
    {ok, "Hello, world"}.  
  
to_html(Req, State) ->  
    Output = io_lib:format("<html><body>~s</body></html>",  
                           [State]),  
    {Output, Req, State}.
```

Must be idempotent

PUT

```
-module(grocery_list_resource).  
-export([init/1,allowed_methods/2]).  
-export([content_types_accepted/2,from_json/2]).  
  
-include_lib("webmachine/include/webmachine.hrl").  
  
init([]) ->  
    {ok, []}.  
  
allowed_methods(Req, State) ->  
    {[ 'PUT' ], Req, State}.  
  
content_types_accepted(Req, State) ->  
    {[ {"application/json", from_json}], Req, State}.  
  
from_json(Req, State) ->  
    %% Create/update logic here
```


DELETE

```
-module(grocery_list_resource).  
-export([init/1,allowed_methods/2]).  
-export([delete_resource/2]).  
  
-include_lib("webmachine/include/webmachine.hrl").  
  
init([]) ->  
    {ok, []}.  
  
allowed_methods(Req, State) ->  
    {[ 'DELETE' ], Req, State}.  
  
delete_resource(Req, State) ->  
    %% Deletion logic here
```

POST

Hmmmm

Problems with POST

- Overloaded semantics
- Create or update?

Creation via POST

- `allowed_methods/2`
- `post_is_create/2`
- `create_path/2`
- `content_types_accepted/2`
- handler function

Update via POST

- `allowed_methods/2`
- `content_types_accepted/2`
- handler function

Lab Time!

- Write a webmachine application to provide a RESTful interface to the cache server.
- URLs should be of the form /cache/<key>
- GETs should fetch cache entries, DELETEs should delete cache entries, and PUTs should store cache entries.